

Non-Life Insurance — Assignment 5

AJ*

Autumn 2021

Question 1

We first check that we have constructed the IBNR triangle correctly, having used the data of Taylor & Ashe (1983):

```
> xtabs(Xij~i+j)
```

i	j	1	2	3	4	5	6	7	8	9	10
1	357848	766940	610542	482940	527326	574398	146342	139950	227229	67948	
2	352118	884021	933894	1183289	445745	320996	527804	266172	425046	0	
3	290507	1001799	926219	1016654	750816	146923	495992	280405	0	0	
4	310608	1108250	776189	1562400	272482	352053	206286	0	0	0	
5	443160	693190	991983	769488	504851	470639	0	0	0	0	
6	396132	937085	847498	805037	705960	0	0	0	0	0	
7	440832	847631	1131398	1063269	0	0	0	0	0	0	
8	359480	1061648	1443370	0	0	0	0	0	0	0	
9	376686	986608	0	0	0	0	0	0	0	0	
10	344014	0	0	0	0	0	0	0	0	0	

We then compute the chain ladder estimates and verify that the model satisfies the *marginal totals property*:

```
## rows marginal totals
> (tapply(Xij, i, sum) - tapply(fitted(Orig.CL), i, sum))/tapply(Xij, i, sum)
```

	1	2	3	4	5	6
	-4.559358e-11	-3.852987e-11	-4.286977e-11	-6.457354e-11	-4.379618e-11	-1.384984e-13
	7	8	9	10		
	-1.684500e-14	-1.398041e-14	-1.468754e-14	-1.235170e-14		

```
## columns marginal totals
> (tapply(Xij, j, sum) - tapply(fitted(Orig.CL), j, sum))/tapply(Xij, j, sum)
```

	1	2	3	4	5	6
	-7.105315e-13	-7.610457e-13	-8.601956e-13	-1.093003e-12	-2.736024e-12	-5.004009e-10
	7	8	9	10		
	-6.888954e-11	-4.344418e-13	-6.646440e-13	-1.349224e-14		

*Student number: ∞

Question 2

`row(Orig.fits)` and `col(Orig.fits)` return matrices that have entries corresponding to the row or column index of that entry. `future` is a logical vector which identifies the entries that are below the diagonal of the matrix.

Question 3

The number of negative pseudo-observations is given by:

```
> sum(n.neg)
[1] 117
```


from a total of $55 \times 1000 = 55\,000$ generated pseudo-observations.

Question 4

Before computing the empirical quantiles, the following visualisation is useful:

```
## kernel density estimate
hist(payments, breaks=21, prob=TRUE, xlab = "Payments")
lines(density(payments), col="blue")
curve(dnorm(x, mean(payments), sd(payments)), add=TRUE, col="red")

## add
legend("topright", c("Kernel density", "Normal density"),
      lty=c(1,1), lwd=c(2.5,2.5), col=c("blue", "red"),
      cex = .75)
```



Now, the empirical quantiles are given by:

```
> quantile(payments, c(.5, .75, .9, .95, .99))
      50%      75%      90%      95%      99%
18.93649 20.93534 22.82899 24.09668 26.30068
```

We see that the generated quantiles are close to those reported in the first paragraph of Remark 10.6.1 in MART.

Question 5

```
## Gamma Response Model
```

```
## (a) Fit:
```

```
Orig.gam <- glm(Xij~i+j, family = Gamma(link=log))
```

```
## (b) Extract coefficients:
```

```
coefs <- exp(coef(Orig.gam))
```

```
alpha <- c(1, coefs[2:TT]) * coefs[1]
```

```
beta <- c(1, coefs[(TT+1):(2*TT-1)])
```

```
names(alpha) <- paste0("row",1:10)
```

```
names(beta)  <- paste0("col",1:10)
```

```
## Compute reserves:
```

```
Orig.fits <- alpha %o% beta; round(Orig.fits)
```

```
future <- row(Orig.fits) + col(Orig.fits) - 1 > TT
```

```
Orig.reserve <- sum(Orig.fits[future])
```

```
## (c) Pearson residuals
```

```
Prs.resid <- (Xij - fitted(Orig.gam)) / fitted(Orig.gam)
```

```
p <- 2*TT-1; phi.P <- sum(Prs.resid^2)/(n-p)
```

```
Adj.Pr.resid <- Prs.resid * sqrt(n/(n-p))
```

continued on next page.

```

## (d) Adjusted bootstrap:
set.seed(174297)
nBoot <- 1000; payments.gam <- reserves.gam <- n.neg <- numeric(nBoot)
for (boots in 1:nBoot){
  Ps.Xij <- sample(Adj.Prs.resid, n, replace=TRUE)
  Ps.Xij <- Ps.Xij * fitted(Orig.gam) + fitted(Orig.gam)
  number.neg <- sum(Ps.Xij<0)
  Ps.Xij <- pmax(Ps.Xij, 0.01)
  Ps.gam <- glm(Ps.Xij~i+j, family = Gamma(link=log))
  coefs <- exp(as.numeric(coef(Ps.gam)))
  Ps.alpha <- c(1, coefs[2:TT]) * coefs[1]
  Ps.beta <- c(1, coefs[(TT+1):(2*TT-1)])
  Ps.fits <- Ps.alpha %o% Ps.beta
  Ps.reserve <- sum(Ps.fits[future])
  h <- length(Ps.fits[future])
  Ps.payments <- rgamma(h, shape=(1/phi.P), rate=(1/(Ps.fits[future]*phi.P)))
  Ps.totpayment <- sum(Ps.payments)
  reserves.gam[boots] <- Ps.reserve
  payments.gam[boots] <- Ps.totpayment
  n.neg[boots] <- number.neg
}

## (e) Statistics (in millions)
pp <- payments-mean(payments)
skew <- mean(pp^3)/sd(payments)^3
kurt <- mean(pp^4)/mean(pp^2)^2 - 3

payments.gam <- payments.gam/1e6
pp.gam <- payments.gam - mean(payments.gam)
skew.gam <- mean(pp.gam^3)/sd(payments.gam)^3
kurt.gam <- mean(pp.gam^4)/mean(pp.gam^2)^2 - 3


## comparison
> rbind(c(mean(payments), sd(payments), skew, kurt),
+       c(mean(payments.gam), sd(payments.gam), skew.gam, kurt.gam))
      [,1]      [,2]      [,3]      [,4]
[1,] 19.00545 2.966485 0.2107881 0.02943967
[2,] 18.31776 2.838490 0.5545964 0.54901357

continued on next page.

```

```
## 5e. continued:
## kernel density estimate
hist(payments.gam, breaks=21, prob=TRUE, xlab = "Payments",
     main = "Histogram of Payments - Gamma Response")
lines(density(payments.gam), col="royalblue")
curve(dnorm(x, mean(payments.gam), sd(payments.gam)), add=TRUE, col="red")

## add
legend("topright", c("Kernel density","Normal density"),
     lty=c(1,1), lwd=c(2.5,2.5), col=c("royalblue" ,"red"),
     cex = .75)
```



Now, the new empirical quantiles are given by:

```
> quantile(payments.gam, c(.5, .75, .9, .95, .99))
      50%      75%      90%      95%      99%
18.01098 20.04397 21.99663 23.41159 26.27565
```

For ease of reference, recall our earlier quantiles:

```
> quantile(payments, c(.5, .75, .9, .95, .99))
      50%      75%      90%      95%      99%
18.93649 20.93534 22.82899 24.09668 26.30068
```

The notable differences are that the mean of the Gamma response is lower and this model exhibits far more skewness and kurtosis.

Question 6

Note: to simplify indexing, this implementation creates a run-off triangle from the data with zeros for future observations, stored as a `matrix` object.

```
Xij <- as.matrix(xtabs(Xij~i+j))

alpha.dm <- numeric(TT)
for (k in 1:TT){
  alpha.dm[k] <- sum(Xij[k,]/beta) / rowSums(!future)[k]
}

> round(alpha - alpha.dm, 3)
  row1  row2  row3  row4  row5  row6  row7  row8  row9 row10
-0.915  0.721  1.009  0.004 -0.316 -0.146 -0.023  0.221  0.018  0.000

beta.dm <- numeric(TT)
for (z in 1:TT){
  beta.dm[z] <- sum(Xij[,z]/alpha) / colSums(!future)[z]
}

> round(beta - beta.dm, 3)
 col1  col2  col3  col4  col5  col6  col7  col8  col9 col10
    0    0    0    0    0    0    0    0    0    0
```

Therefore, it is indeed the case that the parameters `alpha` and `beta` extracted from `Orig.gam` satisfy the *Direct-Method* equations.

Question 7

(a)

```
> Xij[i==TT]
[1] 344014
```

```
> Xij.1[ii==TT]
[1] 344014      0      0      0      0      0      0      0      0      0
```

Yes, the output looks as we'd expect as our goal was to pad the final row with nine zeros which is what we have achieved.

(b)

```
> coef(CL); coef(Orig.CL)
(Intercept)      ii2      ii3      ii4      ii5      ii6
12.506404677  0.331272153  0.321118578  0.305960003  0.219316314  0.270077015
      ii7      ii8      ii9      ii10      jj2      jj3
0.372208424  0.553333059  0.368934194  0.242032956  0.912526274  0.958830628
      jj4      jj5      jj6      jj7      jj8      jj9
1.025997003  0.435276183  0.080056547 -0.006381469 -0.394452205  0.009378211
      jj10
-1.379906692

(Intercept)      i2      i3      i4      i5      i6
12.506404677  0.331272153  0.321118578  0.305960003  0.219316314  0.270077015
      i7      i8      i9      i10      j2      j3
0.372208424  0.553333059  0.368934194  0.242032956  0.912526274  0.958830628
      j4      j5      j6      j7      j8      j9
1.025997003  0.435276183  0.080056547 -0.006381469 -0.394452205  0.009378211
      j10
-1.379906692

> CL$deviance; Orig.CL$deviance
[1] 1903014
[1] 1903014
```

Clearly, weighting out unobserved cells produces the same fitting results.

Question 8

The equivalent methods can be separated as follows:

Mean-deviance estimate $\{1, 3, 5\}$

Pearson estimate $\{2, 4, 6\}$.

Question 9

```
## Chain Ladder (CL) vs. Exposure Model (EE)
```

```
CL <- glm(Xij ~ fi + fj, quasipoisson)
EE <- glm(Xij ~ offset(log(Expo)) + fj, quasipoisson)
phi <- CL$deviance / CL$df.residual
Delta.Dev.Sc <- (EE$deviance - CL$deviance)/phi
Delta.df <- EE$df.residual - CL$df.residual
reject <- Delta.Dev.Sc > qchisq(.95, Delta.df)

> cat("The exposure model", ifelse(reject, "is", "is not"), "rejected",
+     "since the scaled deviance gained by CL is\n",
+     round(Delta.Dev.Sc,1), "with", Delta.df, "extra parameters.\n")
```

The exposure model is rejected since the scaled deviance gained by CL is 30 with 7 extra parameters.

Question 10

```
## consider a new model
CL.off <- glm(Xij~offset(log(Expo))+fi+fj, quasipoisson)

## (a) Fitted value cell (1,1)
> exp(coef(CL))[1]
149.206
> exp(coef(CL.off))[1] * ee[1]
149.206

## (b) Fitted value cell (2,1)
> exp(coef(CL))[1] * exp(coef(CL))[2]
154.8901
> exp(coef(CL.off))[1] * exp(coef(CL.off))[2] * ee[2]
154.8901
```

Clearly, the fitted values are the same for both models.

Question 11

Please note: in this question, all χ_k^2 critical values are taken at the **95th** percentile.

(a)

```
> anova(EE, EE.adj)
```

Analysis of Deviance Table

Model 1: $X_{ij} \sim \text{offset}(\log(\text{Expo})) + f_j$

Model 2: $X_{ij} \sim \text{offset}(\log(\text{Expo})) + i.\text{is}.3 + f_j$

	Resid. Df	Resid. Dev	Df	Deviance
1	28	82.885		
2	27	35.011	1	47.874

Recall out estimate of $\phi = 1.626585$ from question 9.

The decrease in *scaled* deviance is:

$$\frac{47.874}{1.626585} = 29.43221$$

Comparing this with the critical value of $\chi_1^2 = 3.841$, we reject the null and go with the finer model EE.adj.

(b)

```
> anova(EE.adj, CL)
```

Analysis of Deviance Table

Model 1: $X_{ij} \sim \text{offset}(\log(\text{Expo})) + i.\text{is}.3 + f_j$

Model 2: $X_{ij} \sim f_i + f_j$

	Resid. Df	Resid. Dev	Df	Deviance
1	27	35.011		
2	21	34.158	6	0.85293

Noting that the critical value of $\chi_6^2 = 12.59159$, we do not even need to compute the change in *scaled* deviance to conclude that the observed change is not significant. Therefore, the EE.adj model is preferred.

(c) EE.adj.

Question 12

```
## models for adjusted exposure
EE.typo <- glm(Xij ~ offset(log(Expo1)) + fj, quasipoisson)
EE.adj.typo <- glm(Xij~offset(log(Expo1))+i.is.3+fj, quasipoisson)
```

```
> anova(EE.typo, EE.adj.typo)
```

Analysis of Deviance Table

Model 1: Xij ~ offset(log(Expo1)) + fj

Model 2: Xij ~ offset(log(Expo1)) + i.is.3 + fj

	Resid. Df	Resid. Dev	Df	Deviance
1	28	37.932		
2	27	35.011	1	2.9203

Changing the exposure for year 3, we now see that the change in deviance is not significant, so the change in *scaled* deviance certainly is not. Therefore, the preferred model is now `EE.typo`.

```
> anova(EE.typo, CL)
```

Analysis of Deviance Table

Model 1: Xij ~ offset(log(Expo1)) + fj

Model 2: Xij ~ fi + fj

	Resid. Df	Resid. Dev	Df	Deviance
1	28	37.932		
2	21	34.158	7	3.7733

We have $\chi^2_7 = 14.06714$, and by the above-mentioned logic we see that `EE.typo` is the preferred model *i.e.* the model with column parameters and offset by exposure.

Question 13

```
M <- ee * alpha[1] / ee[1]

> CL.fits <- alpha %o% beta; round(CL.fits, 2)
      fj2  fj3  fj4  fj5  fj6  fj7 fj8
149.21 40.24 10.03 5.20 3.01 1.82 0.49  0
fi2 154.89 41.78 10.41 5.40 3.13 1.89 0.51  0
fi3 188.01 50.71 12.63 6.55 3.80 2.29 0.62  0
fi4 201.15 54.26 13.52 7.01 4.06 2.45 0.66  0
fi5 196.10 52.89 13.18 6.84 3.96 2.39 0.64  0
fi6 271.52 73.24 18.24 9.46 5.48 3.31 0.89  0
fi7 233.12 62.88 15.66 8.13 4.71 2.84 0.77  0
fi8 221.00 59.61 14.85 7.70 4.46 2.69 0.73  0

> BF.fits <- M %o% beta; round(BF.fits, 2)
      fj2  fj3  fj4  fj5  fj6  fj7 fj8
[1,] 149.21 40.24 10.03 5.20 3.01 1.82 0.49  0
[2,] 153.35 41.36 10.30 5.35 3.10 1.87 0.50  0
[3,] 160.50 43.29 10.78 5.59 3.24 1.96 0.53  0
[4,] 167.21 45.10 11.23 5.83 3.38 2.04 0.55  0
[5,] 178.84 48.24 12.02 6.23 3.61 2.18 0.59  0
[6,] 186.92 50.42 12.56 6.52 3.77 2.28 0.61  0
[7,] 190.86 51.48 12.82 6.65 3.85 2.33 0.63  0
[8,] 188.82 50.93 12.69 6.58 3.81 2.30 0.62  0
```

Question 14

```
## (a)
future <- row(CL.fits) + col(CL.fits) - 1 > TT
> CL.reserve <- sum(CL.fits[future]); CL.reserve
[1] 152.0312
> BF.reserve <- sum(BF.fits[future]); BF.reserve
[1] 125.9025

## (b)
> CL.retro <- sum(CL.fits[!future]); CL.retro
[1] 2121
> BF.retro <- sum(BF.fits[!future]); BF.retro
[1] 1810.341

## (c)
> sum(Xij)
[1] 2121
```

Indeed, `CL.retro` is equal to `sum(Xij)` but `BF.retro` is not. This is because the Chain Ladder method satisfies the *marginal totals property*, whereas no such condition applies to the Bornhuetter-Ferguson method.